

Importing packages

```
In [1]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.svm import SVC
from sklearn.naive_bayes import GaussianNB #naive bayes
from sklearn import tree #decision tree
from sklearn.metrics import accuracy_score
from sklearn.metrics import roc_auc_score # for checking gini score
import catboost
from catboost import CatBoostClassifier
```

Loading the data

```
In [2]: df = pd.read_csv(r'C:/Users/aksha_0iafvb0/OneDrive/Desktop/creditcardclg.csv')
```

```
In [3]: # prints first 5 rows of the table
df.head()
```

Out[3]:

	Time	V1	V2	V3	V4	V5	V6	V7	V8	V9	...	V21	V22	V23	V24	V25
0	0.0	-1.359807	-0.072781	2.536347	1.378155	-0.338321	0.462388	0.239599	0.098698	0.363787	...	-0.018307	0.277838	-0.110474	0.066928	0.1281
1	0.0	1.191857	0.266151	0.166480	0.448154	0.060018	-0.082361	-0.078803	0.085102	-0.255425	...	-0.225775	-0.638672	0.101288	-0.339846	0.1671
2	1.0	-1.358354	-1.340163	1.773209	0.379780	-0.503198	1.800499	0.791461	0.247676	-1.514654	...	0.247998	0.771679	0.909412	-0.689281	-0.3271
3	1.0	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203	0.237609	0.377436	-1.387024	...	-0.108300	0.005274	-0.190321	-1.175575	0.6471
4	2.0	-1.158233	0.877737	1.548718	0.403034	-0.407193	0.095921	0.592941	-0.270533	0.817739	...	-0.009431	0.798278	-0.137458	0.141267	-0.2061

5 rows × 31 columns

```
In [4]: # prints last 5 rows of the table
df.tail()
```

Out[4]:

	Time	V1	V2	V3	V4	V5	V6	V7	V8	V9	...	V21	V22	V23	V24	V25
284802	172786.0	-11.881118	10.071785	-9.834783	-2.066656	-5.364473	-2.606837	-4.918215	7.305334	1.914428	...	0.213454	0.111864	1.014480	-0.50934	
284803	172787.0	-0.732789	-0.055080	2.035030	-0.738589	0.868229	1.058415	0.024330	0.294869	0.584800	...	0.214205	0.924384	0.012463	-1.01622	
284804	172788.0	1.919565	-0.301254	-3.249640	-0.557828	2.630515	3.031260	-0.296827	0.708417	0.432454	...	0.232045	0.578229	-0.037501	0.64013	
284805	172788.0	-0.240440	0.530483	0.702510	0.689799	-0.377961	0.623708	-0.686180	0.679145	0.392087	...	0.265245	0.800049	-0.163298	0.12320	
284806	172792.0	-0.533413	-0.189733	0.703337	-0.506271	-0.012546	-0.649617	1.577006	-0.414650	0.486180	...	0.261057	0.643078	0.376777	0.00879	

5 rows × 31 columns

Checking the shape of data

```
In [5]: df.shape
```

Out[5]: (284807, 31)

Checking for null values

```
In [6]: df.isnull().sum()
```

```
Out[6]: Time      0
V1      0
V2      0
V3      0
V4      0
V5      0
V6      0
V7      0
V8      0
V9      0
V10     0
V11     0
V12     0
V13     0
V14     0
V15     0
V16     0
V17     0
V18     0
V19     0
V20     0
V21     0
V22     0
V23     0
V24     0
V25     0
V26     0
V27     0
V28     0
Amount  0
Class   0
dtype: int64
```

```
In [7]: # print full summary
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 284807 entries, 0 to 284806
Data columns (total 31 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   Time        284807 non-null float64
 1   V1          284807 non-null float64
 2   V2          284807 non-null float64
 3   V3          284807 non-null float64
 4   V4          284807 non-null float64
 5   V5          284807 non-null float64
 6   V6          284807 non-null float64
 7   V7          284807 non-null float64
 8   V8          284807 non-null float64
 9   V9          284807 non-null float64
10  V10         284807 non-null float64
11  V11         284807 non-null float64
12  V12         284807 non-null float64
13  V13         284807 non-null float64
14  V14         284807 non-null float64
15  V15         284807 non-null float64
16  V16         284807 non-null float64
17  V17         284807 non-null float64
18  V18         284807 non-null float64
19  V19         284807 non-null float64
20  V20         284807 non-null float64
21  V21         284807 non-null float64
22  V22         284807 non-null float64
23  V23         284807 non-null float64
24  V24         284807 non-null float64
25  V25         284807 non-null float64
26  V26         284807 non-null float64
27  V27         284807 non-null float64
28  V28         284807 non-null float64
29  Amount      284807 non-null float64
30  Class       284807 non-null int64
dtypes: float64(30), int64(1)
memory usage: 67.4 MB
```

```
In [8]: # describe the data
df.describe()
```

```
Out[8]:
```

	Time	V1	V2	V3	V4	V5	V6	V7	V8	V9
count	284807.000000	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05
mean	94813.859575	3.918649e-15	5.682686e-16	-8.761736e-15	2.811118e-15	-1.552103e-15	2.040130e-15	-1.698953e-15	-1.893285e-16	-3.147640e-15
std	47488.145955	1.958696e+00	1.651309e+00	1.516255e+00	1.415869e+00	1.380247e+00	1.332271e+00	1.237094e+00	1.194353e+00	1.098632e+00
min	0.000000	-5.640751e+01	-7.271573e+01	-4.832559e+01	-5.683171e+00	-1.137433e+02	-2.616051e+01	-4.355724e+01	-7.321672e+01	-1.343407e+01
25%	54201.500000	-9.203734e-01	-5.985499e-01	-8.903648e-01	-8.486401e-01	-6.915971e-01	-7.682956e-01	-5.540759e-01	-2.086297e-01	-6.430976e-01
50%	84692.000000	1.810880e-02	6.548556e-02	1.798463e-01	-1.984653e-02	-5.433583e-02	-2.741871e-01	4.010308e-02	2.235804e-02	-5.142873e-02
75%	139320.500000	1.315642e+00	8.037239e-01	1.027196e+00	7.433413e-01	6.119264e-01	3.985649e-01	5.704361e-01	3.273459e-01	5.971390e-01
max	172792.000000	2.454930e+00	2.205773e+01	9.382558e+00	1.687534e+01	3.480167e+01	7.330163e+01	1.205895e+02	2.000721e+01	1.559499e+01

8 rows × 31 columns

Checking number of records of each kind of transaction class (Fraud and Non-Fraud)

```
In [9]: count_classes = pd.value_counts(df['Class'], sort = True)
count_classes.plot(kind = 'bar', rot=0)
plt.title("Transaction class distribution")
plt.xticks(range(2))
plt.xlabel("Class")
plt.ylabel("Frequency")
```

```
Out[9]: Text(0, 0.5, 'Frequency')
```



```
In [10]: df['Class'].nunique()
```

```
Out[10]: 2
```

```
In [11]: df.groupby(['Class'])['Class'].count()
```

```
Out[11]: Class
0      284315
1         492
Name: Class, dtype: int64
```

```
In [12]: (492/(492+284315))*100
```

```
Out[12]: 0.1727485630620034
```

```
In [13]: # The data set is highly imbalanced. Looking at each of the fraud(1) and non-fraud(0) transactions.
```

```
frauds = df[df['Class']==1]
normal = df[df['Class']==0]
```

```
In [14]: frauds.shape
```

```
Out[14]: (492, 31)
```

```
In [15]: normal.shape
```

```
Out[15]: (284315, 31)
```

Checking the amount of money involved in each kind of transaction

```
In [16]: # Fraud transactions
frauds.Amount.describe()
```


```
Out[16]: count    492.000000
mean      122.211321
std       256.683288
min        0.000000
25%        1.000000
50%        9.250000
75%       105.890000
max      2125.870000
Name: Amount, dtype: float64
```

```
In [17]: # Non-fraud transactions
normal.Amount.describe()
```

```
Out[17]: count    284315.000000
mean        88.291022
std        250.105092
min         0.000000
25%         5.650000
50%        22.000000
75%        77.050000
max     25691.160000
Name: Amount, dtype: float64
```


```
In [18]: # Box plot of frauds transactions
sns.boxplot(x=frauds['Amount'])
```

```
Out[18]: <AxesSubplot:xlabel='Amount'>
```



```
In [19]: # box plot of normal transactions
sns.boxplot(x = normal['Amount'])
```

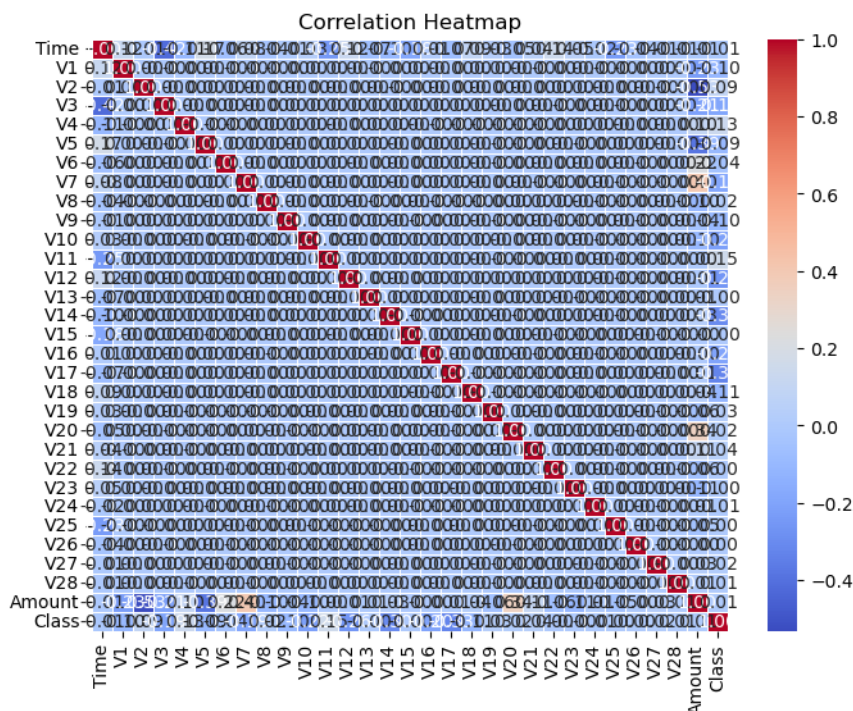
```
Out[19]: <AxesSubplot:xlabel='Amount'>
```



```
In [20]: # heat map of correlation between variables
#sns.heatmap(df, cmap="RdYlGn", annot=True)

correlation_matrix = df.corr()

# Create a heatmap
plt.figure(figsize=(8, 6))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', fmt='.2f', linewidths=.5)
plt.title('Correlation Heatmap')
plt.show()
```



```
In [21]: # separating the dependent and independent variables and stores in x and y variables
x=df[['Time', 'V1', 'V2', 'V3', 'V4', 'V5', 'V6', 'V7', 'V8', 'V9', 'V10',
      'V11', 'V12', 'V13', 'V14', 'V15', 'V16', 'V17', 'V18', 'V19', 'V20',
      'V21', 'V22', 'V23', 'V24', 'V25', 'V26', 'V27', 'V28', 'Amount']]
y=df[['Class']]
```

```
In [22]: y
```

Out[22]:

	Class
0	0
1	0
2	0
3	0
4	0
...	...
284802	0
284803	0
284804	0
284805	0
284806	0

284807 rows × 1 columns

```
In [23]: x
```

Out[23]:

	Time	V1	V2	V3	V4	V5	V6	V7	V8	V9	...	V20	V21	V22	V23
0	0.0	-1.359807	-0.072781	2.536347	1.378155	-0.338321	0.462388	0.239599	0.098698	0.363787	...	0.251412	-0.018307	0.277838	-0.110...
1	0.0	1.191857	0.266151	0.166480	0.448154	0.060018	-0.082361	-0.078803	0.085102	-0.255425	...	-0.069083	-0.225775	-0.638672	0.101...
2	1.0	-1.358354	-1.340163	1.773209	0.379780	-0.503198	1.800499	0.791461	0.247676	-1.514654	...	0.524980	0.247998	0.771679	0.909...
3	1.0	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203	0.237609	0.377436	-1.387024	...	-0.208038	-0.108300	0.005274	-0.190...
4	2.0	-1.158233	0.877737	1.548718	0.403034	-0.407193	0.095921	0.592941	-0.270533	0.817739	...	0.408542	-0.009431	0.798278	-0.137...
...
284802	172786.0	-11.881118	10.071785	-9.834783	-2.066656	-5.364473	-2.606837	-4.918215	7.305334	1.914428	...	1.475829	0.213454	0.111864	1.014...
284803	172787.0	-0.732789	-0.055080	2.035030	-0.738589	0.868229	1.058415	0.024330	0.294869	0.584800	...	0.059616	0.214205	0.924384	0.012...
284804	172788.0	1.919565	-0.301254	-3.249640	-0.557828	2.630515	3.031260	-0.296827	0.708417	0.432454	...	0.001396	0.232045	0.578229	-0.037...
284805	172788.0	-0.240440	0.530483	0.702510	0.689799	-0.377961	0.623708	-0.686180	0.679145	0.392087	...	0.127434	0.265245	0.800049	-0.163...
284806	172792.0	-0.533413	-0.189733	0.703337	-0.506271	-0.012546	-0.649617	1.577006	-0.414650	0.486180	...	0.382948	0.261057	0.643078	0.376...

284807 rows × 30 columns

Spiting the data into 80% training and 20% testing

```
In [24]: # train test split
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.2, random_state=2)
```

```
In [25]: print(x_train.shape,y_train.shape,x_test.shape,y_test.shape)

(227845, 30) (227845, 1) (56962, 30) (56962, 1)
```

```
In [26]: # different model functions
y_train.groupby(['Class'])['Class'].count()
```

Out[26]:

Class	
0	227437
1	408

Name: Class, dtype: int64

```
In [27]: y_test.groupby(['Class'])['Class'].count()
```

Out[27]:

Class	
0	56878
1	84

Name: Class, dtype: int64

Training the credit card fraud detection model

1. Logistic Regression Model

```
In [28]: model = LogisticRegression()
# Fit the model to the data
model.fit(x_train, y_train)
```

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\utils\validation.py:993: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

y = column_or_1d(y, warn=True)

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear_model_logistic.py:814: ConvergenceWarning: lbfgs failed to converge (status=1):

STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression (https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

n_iter_i = _check_optimize_result(

```
Out[28]: LogisticRegression()
```

```
In [29]: # Make predictions
y_pred = model.predict(x_test)
```

```
In [30]: set(y_pred)
```

```
Out[30]: {0, 1}
```

```
In [31]: # convert array into pandas dataframe
d=pd.DataFrame(y_pred)
```

```
In [32]: d
```

```
Out[32]:
```

	0
0	0
1	0
2	0
3	0
4	0
...	...
56957	0
56958	0
56959	0
56960	0
56961	0

56962 rows × 1 columns

```
In [33]: gini=2*roc_auc_score(y_test['Class'],d[0])-1 # formula to predict gini score
```

```
In [34]: gini # prints the gini score
```

```
Out[34]: 0.7492439959210944
```

```
In [35]: # checking the accuracy score predicted by the logistic regression model
accuracy = accuracy_score(y_test, y_pred)
print("Accuracy:", accuracy)
```

Accuracy: 0.9988764439450862

2. Random Forest

```
In [36]: rf_classifier = RandomForestClassifier()
# Train the classifier
rf_classifier.fit(x_train, y_train)
```

C:\Users\aksha_0iafvb0\AppData\Local\Temp\ipykernel_27384\662079852.py:3: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

rf_classifier.fit(x_train, y_train)

```
Out[36]: RandomForestClassifier()
```

```
In [37]: # Make predictions on the test set
y_pred = rf_classifier.predict(x_test)
```

```
In [38]: # convert array into pandas dataframe
d=pd.DataFrame(y_pred)
```

```
In [39]: d # prints the dataframe after converting into pandas from array
```

```
Out[39]:
```

	0
0	0
1	0
2	0
3	0
4	0
...	...
56957	0
56958	0
56959	0
56960	0
56961	0

56962 rows × 1 columns

```
In [40]: gini=2*roc_auc_score(y_test['Class'],d[0])-1 # formula to predict gini score
```

```
In [41]: gini #prints the gini score
```

```
Out[41]: 0.8213406639775358
```

```
In [42]: # Calculate accuracy by Random Forest Classifier
accuracy = accuracy_score(y_test, y_pred)
print("Accuracy:", accuracy)
```

Accuracy: 0.9996488887328394

3. Support vector Machine (SVC)

```
In [43]: svm_classifier = SVC()

# Train the classifier
svm_classifier.fit(x_train, y_train)
```

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\utils\validation.py:993: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().
y = column_or_1d(y, warn=True)

```
Out[43]: SVC()
```

```
In [44]: # Make predictions on the test set
y_pred = svm_classifier.predict(x_test)
```

```
In [49]: set(y_pred)
```

```
Out[49]: {0}
```

```
In [50]: # convert array into pandas dataframe
d=pd.DataFrame(y_pred)
```

```
In [51]: d # prints the dataframe after converting into pandas from array
```

```
Out[51]:
```

	0
0	0
1	0
2	0
3	0
4	0
...	...
56957	0
56958	0
56959	0
56960	0
56961	0

56962 rows × 1 columns


```
In [52]: gini=2*roc_auc_score(y_test['Class'],d[0])-1 # formula to predict gini score
```

```
In [53]: gini #prints the gini score
```

```
Out[53]: 0.0
```

```
In [54]: # Calculate accuracy by svc model
accuracy = accuracy_score(y_test, y_pred)

print("Accuracy:", accuracy)
```

Accuracy: 0.9985253326779256

4. Naive's bayes

```
In [55]: # Build a Gaussian Classifier
model = GaussianNB()

# Model training
model.fit(x_train, y_train)

# Predict Output
predicted = model.predict(x_test)
```

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\utils\validation.py:993: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

```
y = column_or_1d(y, warn=True)
```

```
In [56]: # convert array into pandas dataframe
d=pd.DataFrame(predicted)
```

```
In [57]: d # prints the dataframe after converting into pandas from array
```

```
Out[57]:
```

	0
0	0
1	0
2	0
3	0
4	0
...	...
56957	0
56958	0
56959	0
56960	0
56961	0

56962 rows × 1 columns

```
In [58]: gini=2*roc_auc_score(y_test['Class'],d[0])-1 # formula to predict gini score
```

```
In [59]: gini #prints the gini score
```

```
Out[59]: 0.7073410256539059
```

```
In [60]: #checking the accuracy by naive bayes model
accuracy1 = accuracy_score(y_test, predicted)

print("Accuracy:", accuracy1)
```

Accuracy: 0.9926442189529862

5. Decision Tree Classifier

```
In [61]: # defining decision tree classifier

model = tree.DecisionTreeClassifier(min_samples_split=10)

# train data on new data and new target
model=model.fit(x_train, y_train)
```

```
In [62]: prediction = model.predict(x_test) # assign removed data as input
```

```
In [64]: # convert array into pandas dataframe
d=pd.DataFrame(prediction)
```

```
In [65]: d # prints the dataframe after converting into pandas from array
```

```
Out[65]:
```

	0
0	0
1	0
2	0
3	0
4	0
...	...
56957	0
56958	0
56959	0
56960	0
56961	0

56962 rows × 1 columns

```
In [66]: gini=2*roc_auc_score(y_test['Class'],d[0])-1 # formula to predict gini score
```

```
In [67]: gini #prints the gini score
```

```
Out[67]: 0.7496835331762719
```

```
In [68]: #accuracy score by decision tree classifier
accuracy1 = accuracy_score(y_test, prediction)

print("Accuracy:", accuracy1)
```

Accuracy: 0.9993153330290369

plotting a decision tree of the dataset

```
In [69]: tree.plot_tree(model)
```

```
Out[69]: [Text(0.21912379148827624, 0.9814814814814815, 'X[17] <= -2.769\ngini = 0.004\nsamples = 227845\nvalue = [227437, 408]'),
Text(0.1357234314980794, 0.9444444444444444, 'X[10] <= -1.981\ngini = 0.348\nsamples = 361\nvalue = [81, 280]'),
Text(0.08706786171574904, 0.9074074074074074, 'X[26] <= -0.225\ngini = 0.264\nsamples = 320\nvalue = [50, 270]'),
Text(0.06145966709346991, 0.8703703703703703, 'X[27] <= 1.175\ngini = 0.465\nsamples = 106\nvalue = [39, 67]'),
Text(0.040973111395646605, 0.8333333333333334, 'X[14] <= -2.87\ngini = 0.284\nsamples = 76\nvalue = [13, 63]'),
Text(0.030729833546734954, 0.7962962962962963, 'X[29] <= 1.05\ngini = 0.138\nsamples = 67\nvalue = [5, 62]'),
Text(0.020486555697823303, 0.7592592592592593, 'X[28] <= 0.165\ngini = 0.415\nsamples = 17\nvalue = [5, 12]'),
Text(0.010243277848911651, 0.7222222222222222, 'gini = 0.0\nsamples = 11\nvalue = [0, 11]'),
Text(0.030729833546734954, 0.7222222222222222, 'gini = 0.278\nsamples = 6\nvalue = [5, 1]'),
Text(0.040973111395646605, 0.7592592592592593, 'gini = 0.0\nsamples = 50\nvalue = [0, 50]'),
Text(0.05121638924455826, 0.7962962962962963, 'gini = 0.198\nsamples = 9\nvalue = [8, 1]'),
Text(0.08194622279129321, 0.8333333333333334, 'X[29] <= 94.99\ngini = 0.231\nsamples = 30\nvalue = [26, 4]'),
Text(0.07170294494238157, 0.7962962962962963, 'gini = 0.0\nsamples = 25\nvalue = [25, 0]'),
Text(0.09218950064020487, 0.7962962962962963, 'gini = 0.32\nsamples = 5\nvalue = [1, 4]'),
Text(0.11267605633802817, 0.8703703703703703, 'X[11] <= 1.135\ngini = 0.098\nsamples = 214\nvalue = [11, 203]'),
Text(0.10243277848911651, 0.8333333333333334, 'gini = 0.444\nsamples = 6\nvalue = [4, 2]'),
Text(0.12291933418693982, 0.8333333333333334, 'X[5] <= -22.585\ngini = 0.065\nsamples = 208\nvalue = [7, 201]'),
Text(0.11267605633802817, 0.7962962962962963, 'gini = 0.0\nsamples = 2\nvalue = [2, 0]'),
Text(0.13316261203585147, 0.7962962962962963, 'X[4] <= -0.304\ngini = 0.047\nsamples = 206\nvalue = [5, 201]'),
Text(0.13316261203585147, 0.7962962962962963, 'gini = 0.0\nsamples = 1\nvalue = [1, 0]')]
```

6. Catboost Model

```
In [70]: df
```

Out[70]:

	Time	V1	V2	V3	V4	V5	V6	V7	V8	V9	...	V21	V22	V23	
0	0.0	-1.359807	-0.072781	2.536347	1.378155	-0.338321	0.462388	0.239599	0.098698	0.363787	...	-0.018307	0.277838	-0.110474	0.066
1	0.0	1.191857	0.266151	0.166480	0.448154	0.060018	-0.082361	-0.078803	0.085102	-0.255425	...	-0.225775	-0.638672	0.101288	-0.339
2	1.0	-1.358354	-1.340163	1.773209	0.379780	-0.503198	1.800499	0.791461	0.247676	-1.514654	...	0.247998	0.771679	0.909412	-0.689
3	1.0	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203	0.237609	0.377436	-1.387024	...	-0.108300	0.005274	-0.190321	-1.175
4	2.0	-1.158233	0.877737	1.548718	0.403034	-0.407193	0.095921	0.592941	-0.270533	0.817739	...	-0.009431	0.798278	-0.137458	0.141
...
284802	172786.0	-11.881118	10.071785	-9.834783	-2.066656	-5.364473	-2.606837	-4.918215	7.305334	1.914428	...	0.213454	0.111864	1.014480	-0.509
284803	172787.0	-0.732789	-0.055080	2.035030	-0.738589	0.868229	1.058415	0.024330	0.294869	0.584800	...	0.214205	0.924384	0.012463	-1.016
284804	172788.0	1.919565	-0.301254	-3.249640	-0.557828	2.630515	3.031260	-0.296827	0.708417	0.432454	...	0.232045	0.578229	-0.037501	0.640
284805	172788.0	-0.240440	0.530483	0.702510	0.689799	-0.377961	0.623708	-0.686180	0.679145	0.392087	...	0.265245	0.800049	-0.163298	0.123
284806	172792.0	-0.533413	-0.189733	0.703337	-0.506271	-0.012546	-0.649617	1.577006	-0.414650	0.486180	...	0.261057	0.643078	0.376777	0.008

284807 rows x 31 columns

```
In [71]: x=df[['Time', 'V1', 'V2', 'V3', 'V4', 'V5', 'V6', 'V7', 'V8', 'V9', 'V10',
            'V11', 'V12', 'V13', 'V14', 'V15', 'V16', 'V17', 'V18', 'V19', 'V20',
            'V21', 'V22', 'V23', 'V24', 'V25', 'V26', 'V27', 'V28', 'Amount']]
y=df[['Class']]
```

```
In [72]: # train test split
x_temp, x_test, y_temp, y_test = train_test_split(x, y, test_size=0.2, random_state=2)
```

```
In [73]: # train test split
x_train, x_val, y_train, y_val = train_test_split(x_temp, y_temp, test_size=0.25, random_state=2)
```

```
In [74]: # Create CatBoost classifier
cat_model = CatBoostClassifier( depth=5,
                                learning_rate=0.01,
                                iterations=1500,
                                random_seed=2,
                                loss_function='Logloss',
                                eval_metric='AUC',
                                od_type='Iter',
                                boosting_type='Plain',
                                use_best_model=True,
                                one_hot_max_size=2)
```

```
In [75]: # Train the model
model=cat_model.fit(x_train, y_train,early_stopping_rounds=100,eval_set=(x_val,y_val))
```

0:	test: 0.8568622	best: 0.8568622 (0)	total: 231ms	remaining: 5m 46s
1:	test: 0.9022036	best: 0.9022036 (1)	total: 284ms	remaining: 3m 32s
2:	test: 0.9549925	best: 0.9549925 (2)	total: 334ms	remaining: 2m 46s
3:	test: 0.9639464	best: 0.9639464 (3)	total: 383ms	remaining: 2m 23s
4:	test: 0.9649631	best: 0.9649631 (4)	total: 438ms	remaining: 2m 10s
5:	test: 0.9680959	best: 0.9680959 (5)	total: 493ms	remaining: 2m 2s
6:	test: 0.9677843	best: 0.9680959 (5)	total: 546ms	remaining: 1m 56s
7:	test: 0.9709125	best: 0.9709125 (7)	total: 594ms	remaining: 1m 50s
8:	test: 0.9682731	best: 0.9709125 (7)	total: 644ms	remaining: 1m 46s
9:	test: 0.9662186	best: 0.9709125 (7)	total: 697ms	remaining: 1m 43s
10:	test: 0.9655670	best: 0.9709125 (7)	total: 746ms	remaining: 1m 41s
11:	test: 0.9660417	best: 0.9709125 (7)	total: 798ms	remaining: 1m 38s
12:	test: 0.9621671	best: 0.9709125 (7)	total: 862ms	remaining: 1m 38s
13:	test: 0.9620320	best: 0.9709125 (7)	total: 918ms	remaining: 1m 37s
14:	test: 0.9621649	best: 0.9709125 (7)	total: 978ms	remaining: 1m 36s
15:	test: 0.9605312	best: 0.9709125 (7)	total: 1.04s	remaining: 1m 36s
16:	test: 0.9592117	best: 0.9709125 (7)	total: 1.09s	remaining: 1m 35s
17:	test: 0.9610342	best: 0.9709125 (7)	total: 1.15s	remaining: 1m 34s
18:	test: 0.9624094	best: 0.9709125 (7)	total: 1.2s	remaining: 1m 33s

```
In [76]: # Make predictions
y_pred = model.predict(x_test)
```

```
In [77]: # Evaluate the model
accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy: {accuracy}")

Accuracy: 0.9995611109160493
```

```
In [78]: d=pd.DataFrame(y_pred) # converting the array into pandas dataframe
```

```
In [79]: d.groupby([0])[0].count()
```

```
Out[79]: 0
          56893
          1         69
          Name: 0, dtype: int64
```

```
In [80]: y_test.groupby(['Class'])['Class'].count()
```

```
Out[80]: Class
          0    56878
          1      84
          Name: Class, dtype: int64
```

```
In [81]: gini=2*roc_auc_score(y_test['Class'],d[0])-1 # formula to predict gini score
```

```
In [82]: gini
```

```
Out[82]: 0.761816854453726
```

```
In [ ]:
```